kathara lab

bgp: announcing prefixes with frr

Version	1.0
Author(s)	G. Di Battista, M. Patrignani, M. Pizzonia, F. Ricci, M. Rimondini
E-mail	contact@kathara.org
Web	http://www.kathara.org/
Description	a simple bgp announcement; kathara version of a netkit lab

copyright notice

- All the pages/slides in this presentation, including but not limited to, images, photos, animations, videos, sounds, music, and text (hereby referred to as "material") are protected by copyright.
- This material, with the exception of some multimedia elements licensed by other organizations, is property of the authors and/or organizations appearing in the first slide.
- This material, or its parts, can be reproduced and used for didactical purposes within universities and schools, provided that this happens for non-profit purposes.
- Information contained in this material cannot be used within network design projects or other products of any kind.
- Any other use is prohibited, unless explicitly authorized by the authors on the basis of an explicit agreement.
- The authors assume no responsibility about this material and provide this material "as is", with no implicit or explicit warranty about the correctness and completeness of its contents, which may be subject to changes.
- This copyright notice must always be redistributed together with the material, or its portions.

preconditions

- for this lab we assume you have chosen “kathara/frr” as the default image of your Kathará installation
  - execute “kathara settings”
    - select “choose default image”
    - select “kathara/frr”
    - exit from the settings procedure

announcements and traffic flows

- bgp allows a router to offer connectivity to another router
- "offering connectivity" means "promising the delivery to a specific destination"

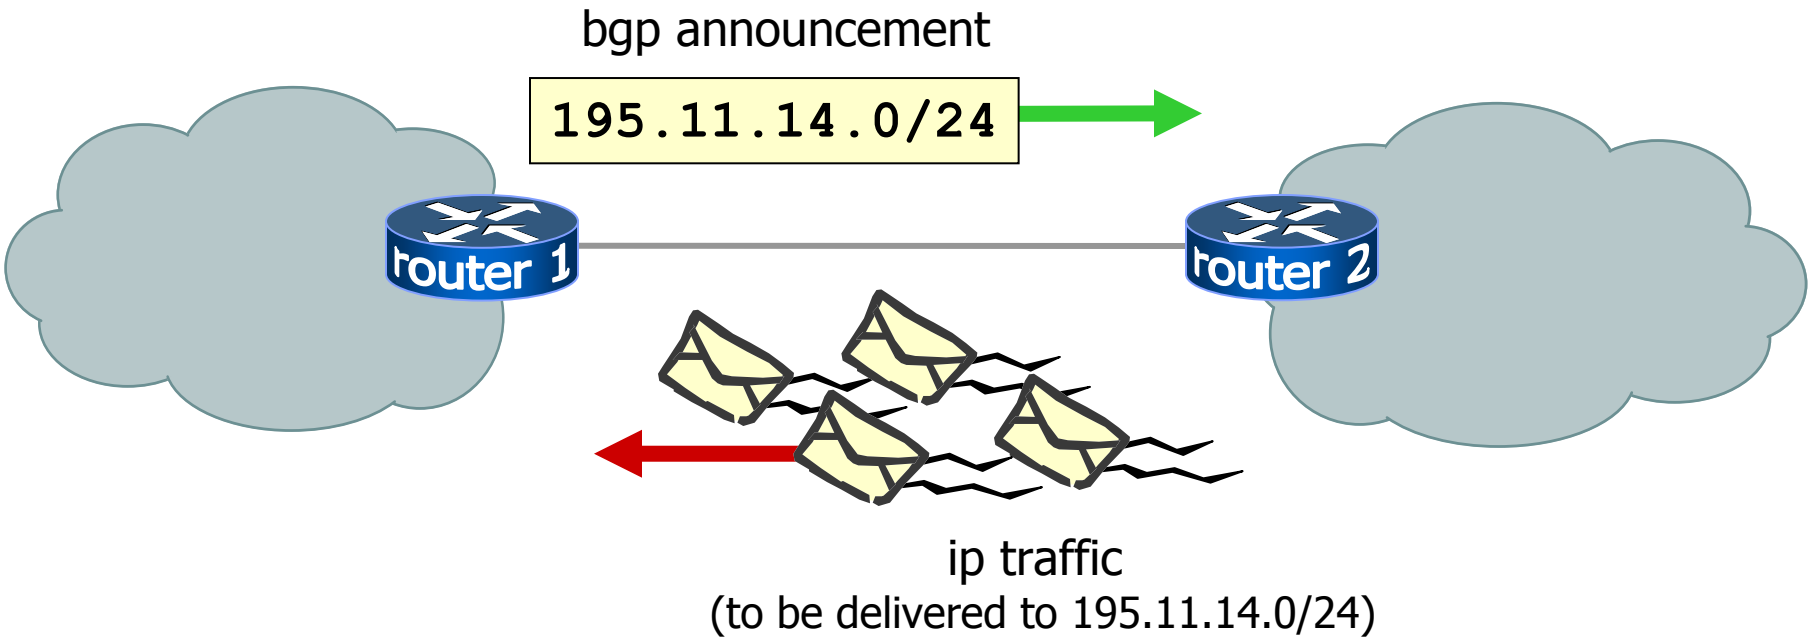

announcement commands

cisco command syntax

```
network <network-ip> mask <network-mask>
```

frr command syntax

```
network <network-ip/network-mask>
```

- the network command performs the following
  - checks whether the current router has the specified prefix in its kernel forwarding table
    - if not, the command has no effect whatsoever
  - for each peer checks whether the BGP configuration has security filters associated with the peer
    - only in the positive case the prefix is announced to the peer

announcement commands

command syntax

```
no bgp network import-check
```

- this command modifies the network statement disabling the check of the existence of the prefix in the kernel forwarding table
  - default value is **`bgp network import-check`**
- by disabling this check you may announce prefixes that
  - may not be local
  - may not even exist (!)
- not needed for this lab

announcement commands

command syntax

```
no bgp ebgp-requires-policy
```

- this command modifies the network statement disabling the check of the existence of incoming and outgoing filters associated with each peer
- if you don't perform this command
  - no prefix will be accepted from a peer if an incoming filter is not set for that peer
  - no prefix will be announced to a peer if an outgoing filter is not set for that peer

announcement commands

- observe that the **`network`** command
  - *does not* inject any route in the kernel forwarding table
  - checks whether the network address matches the netmask; if it does not, the command is automatically replaced in the router configuration; for example:
    - **`network 193.100.0.0/8`**
      is replaced by
      `network 193.0.0.0/8`
    - **`network 1.2.3.4/0`**
      is replaced by
      `network 0.0.0.0/0`

announcement example

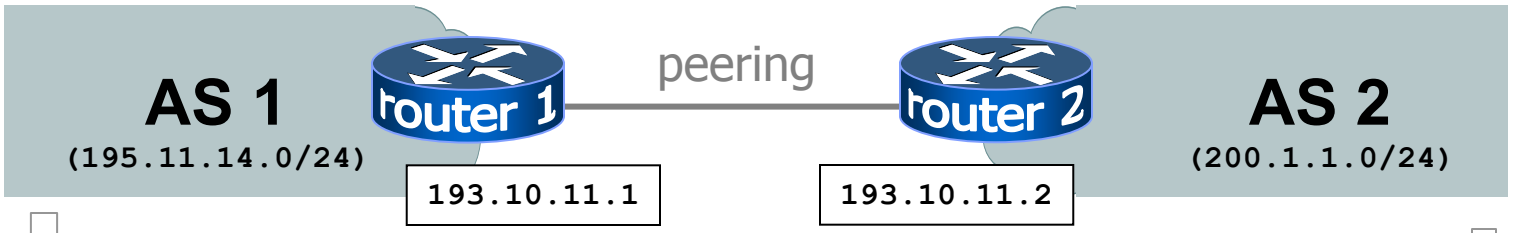


```
! router 1 configuration file
router bgp 1
no bgp ebgp-requires-policy
neighbor 193.10.11.2 remote-as 2
network 195.11.14.0/24
```

```
! router 2 configuration file
router bgp 2
no bgp ebgp-requires-policy
neighbor 193.10.11.1 remote-as 1
network 200.1.1.0/24
```

peering configuration

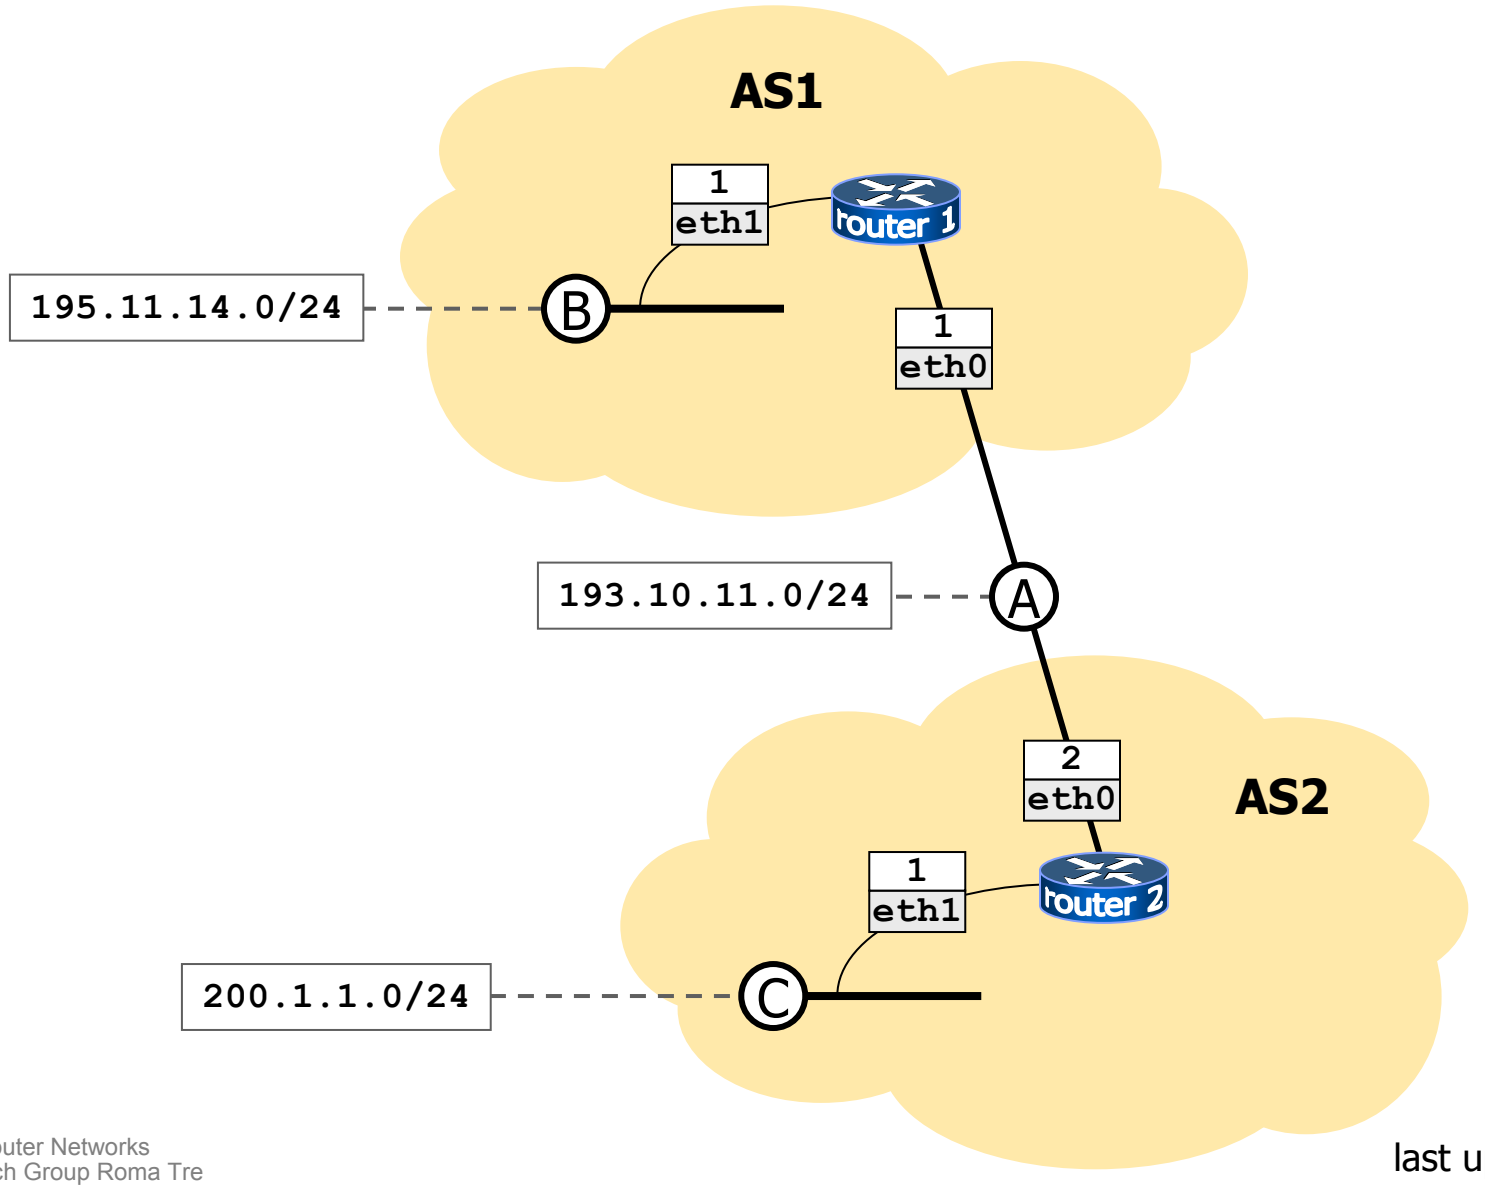

last update: Oct 2021

announcement example

- start the lab

host machine

```
user@localhost:~$ cd kathara-lab_bgp-announcement_frr
user@localhost:~/kathara-lab_bgp-announcement_frr$ kathara lstart
```

- check the zebra routing table

router2

```
root@router2:/# vtysh
...
router2-frr# show ip route
Codes: K - kernel route, C - connected, S - static, R - RIP,
       O - OSPF, I - IS-IS, B - BGP, E - EIGRP, N - NHRP,
       T - Table, v - VNC, V - VNC-Direct, A - Babel, D - SHARP,
       F - PBR, f - OpenFabric,
       > - selected route, * - FIB route, q - queued, r - rejected, b -
backup

C>* 193.10.11.0/24 is directly connected, eth0, 00:03:41
B>* 195.11.14.0/24 [20/0] via 193.10.11.1, eth0, weight 1, 00:03:38
C>* 200.1.1.0/24 is directly connected, eth1, 00:03:41
router2-frr#
```

announcement example

- check the bgpd log file

router2

```
root@router2:/# cat /var/log/frr/frr.log
...
2021/10/23 22:44:06 BGP: 193.10.11.1 sending KEEPALIVE
2021/10/23 22:44:06 BGP: 193.10.11.1 KEEPALIVE rcvd
2021/10/23 22:44:07 BGP: u1:s1 announcing routes upon coalesce timer expiry(1050 ms)
2021/10/23 22:44:07 BGP: u1:s1 send UPDATE w/ attr: nexthop 0.0.0.0, origin i,
mp_nexthop ::, metric 0, path
2021/10/23 22:44:07 BGP: u1:s1 send UPDATE 200.1.1.0/24 IPv4 unicast
2021/10/23 22:44:07 BGP: u1:s1 send UPDATE len 55 numpfx 1
2021/10/23 22:44:07 BGP: u1:s1 193.10.11.1 send UPDATE w/ nexthop 193.10.11.2
2021/10/23 22:44:07 BGP: 193.10.11.1 rcvd UPDATE w/ attr: nexthop 193.10.11.1, origin
i, metric 0, path 1
2021/10/23 22:44:07 BGP: 193.10.11.1 rcvd UPDATE wlen 0 attrlen 28 alen 4
2021/10/23 22:44:07 BGP: 193.10.11.1 rcvd 195.11.14.0/24 IPv4 unicast
2021/10/23 22:44:07 BGP: 193.10.11.1 rcvd UPDATE wlen 0 attrlen 0 alen 0
2021/10/23 22:44:07 BGP: bgp_best_path_select_defer: processing route for IPv4 Unicast
: cnt 0
2021/10/23 22:44:07 BGP: bgp_update_receive: rcvd End-of-RIB for IPv4 Unicast from
193.10.11.1 in vrf default
...
```

announcement example

- check the bgpd log file

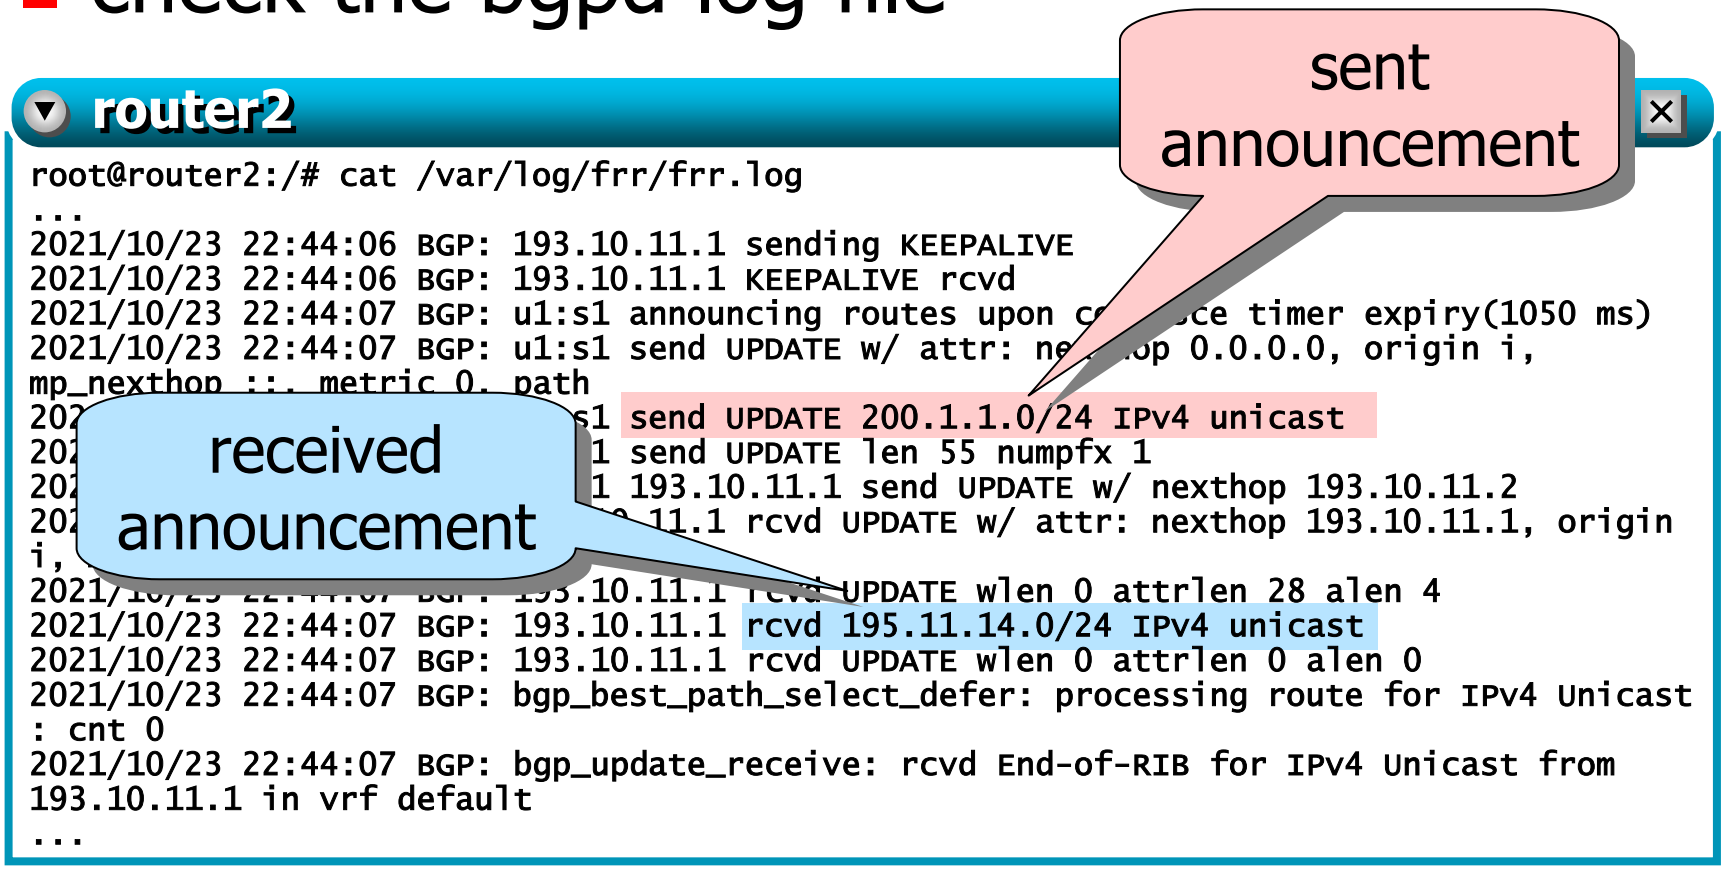

announcement example

- check the vtysh cli (command line interface)
  - type "**`vtysh`**"
  - type "**`show ip bgp neighbors`**"
  - type "**`show ip bgp`**"
  - type "**`show ip bgp 200.1.1.0`**"
- ping "**`200.1.1.0`**"
- terminate the lab

host machine

```
user@localhost:~$ cd kathara-lab_bgp-announcement_frr
user@localhost:~/kathara-lab_bgp-announcement_frr$ kathara lclean
```